

CAPABILITIES OVERVIEW

An AI That Runs Your GoHighLevel

Enterprise-grade automation, built and running on a startup budget. It reads everything, builds entire automations from a sentence, and reaches beyond GHL entirely.

The thesis: most "AI for GHL" tools can only read or suggest. This one **does the work** — constructs and publishes workflows, sends messages, scores contacts, analyzes calls, and automates platforms that have no API at all. All 24/7, for less than a typical SaaS subscription.

A shareable overview. The companion repo contains the full documentation, a runnable read-only demo, and starter code.

Contents

1	What it can do	The capability tour, read & write
2	Built for scale	Dozens of workflows, hundreds of contacts, whole sites, one run
3	vs. GHL's built-in AI	Point features vs. a full operator
4	How it works	The architecture, in plain terms
5	The technical edge	Endpoints, two-token auth, the hard-won gotchas
6	Why a server beats a laptop app	Always-on vs. single-session
7	Getting started	Run the demo, then build your own

How to read this: this PDF is the overview. The companion repo holds the full detail — six documentation files, a read-only demo you run against your own account, and starter code. Tap any section above to jump to it.

1 What it can do

You talk to it like a teammate — "build me a new-lead nurture," "who went cold this week," "enroll everyone who bought" — and it actually does the work inside GHL and beyond, with a confirmation step before anything changes.

Builds automations from a prompt [write](#)

Describe a multi-step workflow in words; it constructs the triggers, branches, messages, waits, and tags — and publishes it live. Multiple flows in one session.

Messaging that sends itself [write](#)

SMS / email / WhatsApp from the right rep's number, with duplicate protection so nobody is ever double-texted.

Pipeline & contact ops [write](#)

Moves deals, updates fields, adds tags — driven by calls, replies, and payments. Bulk actions get a safety pre-flight scan.

AI contact intelligence [write](#)

Scores + summarizes every contact (tier + next action) written back into GHL, plus call-recording sentiment analysis for coaching.

Reaches outside GHL [write](#)

Automates platforms with no API (even behind Cloudflare): create account, wait for activation, enroll, confirm — across two systems, hands-off.

Always-on & conversational

A chat "cockpit" your team uses by text or voice. Background pollers, scheduled jobs, and webhooks run around the clock.

See it on your own account in 60 seconds. The companion repo includes a *read-only* demo: enter your GHL token + location ID and watch it connect, list your pipelines, surface real sales, and generate a sample AI contact review. It only reads — it changes nothing.

2 Built for scale

The single biggest difference in day-to-day use isn't any one action, it's *volume*. One plain-English request fans out into dozens of API calls running in parallel, so jobs that would be a week of agency clicking finish in an afternoon. Three real examples from accounts running today:

Dozens of workflows from one request

Describe a whole automation suite and it builds the entire set, not one flow. A from-scratch account buildout produced a **full library of ~17 workflows** in a single run, new-lead welcome (by lead source), inspection reminders, quote follow-ups, approval drip sequences, build-day reminders, and more, each with real message copy (no placeholder text), proper cadence (immediate / 5 min / 1 hr / 24 hr / N days), merge tags, and stop-on-reply branches. Each workflow is a precise 5-call API chain (create shell → fetch → save steps → register trigger → publish); the system runs many concurrently, tracks which succeeded, and resumes any that failed *without creating duplicates*.

Audit hundreds of contacts, with a personalized next step for each

It paginates every contact and opportunity in the account and, for each one, reads the full history, fields, tags, conversations, even call-recording sentiment, then writes back a **tier** (Champion / Strong / Solid / Pass), a one-line summary, and the recommended next action, directly into GHL custom fields. Roughly ten contacts are processed in parallel. From there it can draft a contact-specific follow-up email for *each* person and queue them for one-click send. This bulk, account-wide review with structured write-back is the headline thing GHL's own AI cannot do.

Build and publish a whole website at once

It stands up and publishes an entire site, not a single page. One real buildout published a **17-page local-service site** (home, service pages, eight city landing pages, plus privacy / terms / sitemap) in one automated pass, then ran SEO, navigation, and logo fixes across every page together. Because GHL's page builder has no public API, this drives a real headless browser that logs in, opens each page, clicks publish, and verifies the result.

The pattern: one sentence from you sets off a fan-out of concurrent, idempotent, self-checking work, an afternoon of agency labor, hands-off, with a human approval step in front of anything that changes data.

3 vs. GHL's built-in AI

GoHighLevel has invested heavily in AI, and credit where it's due, it's a broad suite (marketed as the "AI Employee"). It includes **Conversation AI** (a trainable chatbot for inbound chats, FAQs, and appointment booking), **Voice AI** agents that answer inbound and outbound calls, capture caller details, book appointments, and escalate to a live rep, plus a browser-based **Voice AI chat widget**, **Content AI** (generates social posts, emails, headlines, and blogs inside the builders), **Reviews AI** (suggests or auto-posts review replies), the **Workflow AI Assistant** (a chat helper inside the workflow builder that explains, optimizes, and adds steps, and can draft SMS/email copy and custom code), and **Funnel & Website AI / AI Studio** (generates page layout, copy, and visuals from a prompt, URL, or image). You can also drop OpenAI/ChatGPT into a workflow as an **AI action**.


That's a lot of surface area, and yet it all shares one ceiling: each is an *in-app assistant scoped to a single screen, driven by a human, working one item at a time*. The Workflow AI Assistant builds a single workflow at a time (with a daily usage cap) and helps you in the builder, it doesn't construct and publish an entire library on its own. Funnel & Website AI accelerates one page inside the builder; it doesn't publish a 17-page site and fix SEO across every page in one unattended run. And none of them audit hundreds of contacts with structured write-back, coach from call recordings, reach platforms outside GHL, or run multi-step jobs unattended, because none of them have direct API access. They only have the buttons on their own screens.

Capability	GHL's native AI suite	This system
Chat / answer questions about your data	Yes, in-app, per screen	Yes, across the whole account
Generate content & draft copy	Yes (Content AI)	Yes
Help build a workflow	Yes, one at a time, assisted	Yes, builds & publishes dozens per run
Generate a web page	Yes, one page, in-builder	Yes, whole sites + site-wide SEO
Audit hundreds of contacts → write back tier + next step	No	Yes
Bulk personalized outbound, with a safety pre-flight	No	Yes
Analyze call recordings for sentiment / coaching	No	Yes
Reach platforms outside GHL (even no-API ones)	No	Yes
Direct API / endpoint access (not just UI buttons)	No	Yes
Run unattended across multi-step jobs	Limited (answers calls / chats)	Yes, full orchestration

The unlock is the endpoints. GHL's native AI can only touch what its own screens expose, one item at a time. This system talks *directly* to GHL's APIs, including the internal admin API the GHL web app itself uses to construct workflows and funnels. That's the gap between an assistant that helps you click the buttons on a page and one wired into the same plumbing the platform runs on, plus a real browser for anything that has no API at all.

4 How it works

Five parts on one small always-on server: a web "cockpit" you chat with, the AI brain (Claude), modular tool servers that give the AI safe abilities, the CRM/external services, and a local database that remembers everything. Writes route through a human confirmation step.



Two design choices that matter. (1) The AI runs as a CLI subprocess, so it bills against a flat-rate subscription instead of metered per-token API calls — a big cost difference at scale. (2) Every data-changing action is *queued for human approval* before it runs, which is what makes letting an AI operate a live CRM safe.

5 The technical edge

The companion repo includes the full GoHighLevel API knowledge base, months of reverse-engineering distilled into one reference: every endpoint, the **two-token auth model**, workflow step shapes, and the *gotchas* that silently break things (the ones that cost weeks to discover). A few that matter most:

- **Two tokens, two surfaces.** A stable public API token handles contacts, tags, pipelines, messaging, payments, and calendars; a short-lived admin JWT (≈1 hr) unlocks the internal workflow and funnel APIs the GHL app uses behind the scenes. Knowing which is which is half the battle.
- **The 5-call workflow chain.** Creating one live workflow is an exact sequence, create shell → re-fetch → save steps → register trigger (`triggersChanged: true` is mandatory, fails silently without it) → publish. Miss a step or reuse a stale object and you get a cryptic 422.
- **Destructive defaults.** A `PUT /contacts/{id}` that includes a `tags` field wipes every tag on the contact. Tags must go through the additive `POST /contacts/{id}/tags`. One missed detail like this can silently mangle an account.
- **Version + rate-limit quirks.** Workflows and funnels speak API version `2021-04-15`; everything else speaks `2021-07-28`. Limits sit near ~100 requests/min per location, so bulk jobs are throttled and batched automatically.
- **Beyond GHL.** For platforms with no API at all (even ones behind Cloudflare), a documented browser-automation method drives a real Chrome session: log in, navigate, act, verify, the same way the website-publishing and cross-platform enrollment features work.

It's a shared, living knowledge base. Multiple AI agents contribute verified endpoints and gotchas back to it, so the reference gets sharper with every buildout instead of going stale.

6 Why a server beats a laptop app

	Desktop app	Always-on server
Runs while you're asleep	No	Yes
Background pollers / webhooks	No	Yes
Long tasks survive closing the tab	No	Yes
Whole team, one cockpit	No	Yes

7 Getting started

- **Run the demo** against your own GHL (read-only) — proof in 60 seconds.
- **Read the capability tour** (docs/01-what-it-can-do.md) for use cases.
- **Hand the rebuild guide** (docs/04-rebuild-guide.md) + the knowledge base to your AI assistant and build your own, one stage at a time.
- **Stand up the server** (docs/06-full-server-setup.md) when you want it running 24/7.

Everything described here is real and running today — on a startup budget. The repo gives you the documentation, the demo, and the starter code to build the same thing.