

CAPABILITIES OVERVIEW

An AI That Runs Your GoHighLevel

Enterprise-grade automation, built and running on a startup budget. It reads everything, builds entire automations from a sentence, and reaches beyond GHL entirely.

The thesis: most "AI for GHL" tools can only read or suggest. This one **does the work** — constructs and publishes workflows, sends messages, scores contacts, analyzes calls, and automates platforms that have no API at all. All 24/7, for less than a typical SaaS subscription.

A shareable overview. The companion repo contains the full documentation, a runnable read-only demo, and starter code.

What it can do

You talk to it like a teammate — "build me a new-lead nurture," "who went cold this week," "enroll everyone who bought" — and it actually does the work inside GHL and beyond, with a confirmation step before anything changes.

Builds automations from a prompt

[write](#)

Describe a multi-step workflow in words; it constructs the triggers, branches, messages, waits, and tags — and publishes it live. Multiple flows in one session.

Messaging that sends itself

[write](#)

SMS / email / WhatsApp from the right rep's number, with duplicate protection so nobody is ever double-texted.

Pipeline & contact ops

[write](#)

Moves deals, updates fields, adds tags — driven by calls, replies, and payments. Bulk actions get a safety pre-flight scan.

AI contact intelligence

[write](#)

Scores + summarizes every contact (tier + next action) written back into GHL, plus call-recording sentiment analysis for coaching.

Reaches outside GHL

[write](#)

Automates platforms with no API (even behind Cloudflare): create account, wait for activation, enroll, confirm — across two systems, hands-off.

Always-on & conversational

A chat "cockpit" your team uses by text or voice. Background pollers, scheduled jobs, and webhooks run around the clock.


See it on your own account in 60 seconds. The companion repo includes a *read-only* demo: enter your GHL token + location ID and watch it connect, list your pipelines, surface real sales, and generate a sample AI contact review. It only reads — it changes nothing.

Why it's different

Capability	Typical AI/GHL tools	This system
Answer questions about your data	Yes	Yes
Build & publish real workflows	No	Yes
Send messages safely (no double-texts)	Limited	Yes
Score + summarize every contact	Rare	Yes
Automate platforms with no API	No	Yes
Run 24/7 for the whole team	Varies	Yes

How it works

Five parts on one small always-on server: a web "cockpit" you chat with, the AI brain (Claude), modular tool servers that give the AI safe abilities, the CRM/external services, and a local database that remembers everything. Writes route through a human confirmation step.



Two design choices that matter. (1) The AI runs as a CLI subprocess, so it bills against a flat-rate subscription instead of metered per-token API calls — a big cost difference at scale. (2) Every data-changing action is *queued for human approval* before it runs, which is what makes letting an AI operate a live CRM safe.

The technical edge

The companion repo includes the full GoHighLevel API knowledge base — months of reverse-engineering distilled into one reference: every endpoint, the two-token auth model (stable public API +

short-lived admin API for workflow internals), workflow step shapes, and the *gotchas* that silently break things (the ones that cost weeks to discover). It also documents the browser-automation method for reaching any platform that has no API.

Why a server beats a laptop app

	Desktop app	Always-on server
Runs while you're asleep	No	Yes
Background pollers / webhooks	No	Yes
Long tasks survive closing the tab	No	Yes
Whole team, one cockpit	No	Yes

Getting started

- **Run the demo** against your own GHL (read-only) — proof in 60 seconds.
- **Read the capability tour** (docs/01-what-it-can-do.md) for use cases.
- **Hand the rebuild guide** (docs/04-rebuild-guide.md) + the knowledge base to your AI assistant and build your own, one stage at a time.
- **Stand up the server** (docs/06-full-server-setup.md) when you want it running 24/7.

Everything described here is real and running today — on a startup budget. The repo gives you the documentation, the demo, and the starter code to build the same thing.